Part I
Reinforcement Learning

1 Purpose

The topic for this week is Reinforcement Learning (RL). We will see how reinforcement learning tasks can be formally defined, compare different models of optimality, and discuss the need for exploration and the problem of delayed rewards. We will present a number of RL algorithms including Temporal Difference learning (TD-learning), Q-learning, policy gradients, actor-critic. We will see how RL can be applied to games like backgammon, and how deep Q-learning or A3C can learn to play Atari games from raw pixels.

1.1 Activity Instructions

Complete the exercises (both non-coding and coding) in Ed to help you check your understanding and prepare for your assessments.

2 Week 5: Overview

The topic for this week is Reinforcement Learning (RL). We will see how reinforcement learning tasks can be formally defined, compare different models of optimality, and discuss the need for exploration and the roblem of delayed rewards. We will present a number of RL algorithms including Temporal Difference learning (TD-learning), Q-learning, policy gradients, actor-critic. We will see how RL can be applied to games like backgammon, and how deep Q-learning or A3C can learn to play Atari games from raw pixels.

2.1 Weekly learning outcomes

By the end of this week, you will be able to:

1. explain the difference between supervised learning, reinforcement learning, and unsupervised learning
2. describe the formal definition of a reinforcement learning task
3. identify different models of optimality

4. explain the need for exploration in reinforcement learning
5. describe reinforcement learning algorithms, including TD-learning, Q-learning, policy gradients, and actor-critic
6. apply Q-learning to simple tasks.

2.2 Reinforcement Learning

We have previously discussed supervised learning, where pairs of inputs and target outputs are provided and the system must learn to predict the correct output for each input.

There are many situations where we instead want to train a system to perform certain actions in an environment, in order to maximise a reward function. These situations include, for example, playing a video game, allocating mobile phone channels or other resources dynamically, driving a car or flying a helicopter.

Supervised learning can sometimes be used for this purpose, if we construct a training set of situation-action pairs (for example, a database of game positions and the move chosen by a human expert, or sensor readings from a motor car and the steering direction chosen by a human driver). This process is sometimes called Behavioral Cloning.

However, it is often better if the system can learn by purely by self-play, without the need for training data from a human expert.

2.3 Reinforcement Learning Framework

Reinforcement Learning (RL) can be formalised in terms of an Agent interacting with its Environment.

The Environment includes a set of S states and a set of A actions. At each time step t, the agent is in some state s_t. It must choose an action a_t, whereupon it goes into state $s_{t+1} = \delta(s_t, a_t)$ and receives reward $r_t = R(s_t, a_t)$

The Agent chooses its actions according to some policy $\pi : S \to A$

The aim of Reinforcement Learning is to find an optimal policy which maximises the cumulative reward.

In some cases, the Environment may be probabilistic or stochastic meaning that the transitions and/or rewards are not fully determined, but instead occur randomly according to some probability distribution:

$$\delta(s_{t=1} = s|s_t, a_t), R(r_t = r|s_t, a_t)$$

The Agent can also employ a probabilistic policy, with its actions chosen according to a probability distribution:

$$\pi(a_t = a|s_t)$$

2.4 Optional Video

insert video

2.5 Probabilistic Policies

The main benefit of a probabilistic policy is that it forces the agent to explore its environment more comprehensively while it is learning.

However, probabilistic policies may have additional benefits in other contexts, such as partially observable or multiplayer environments.

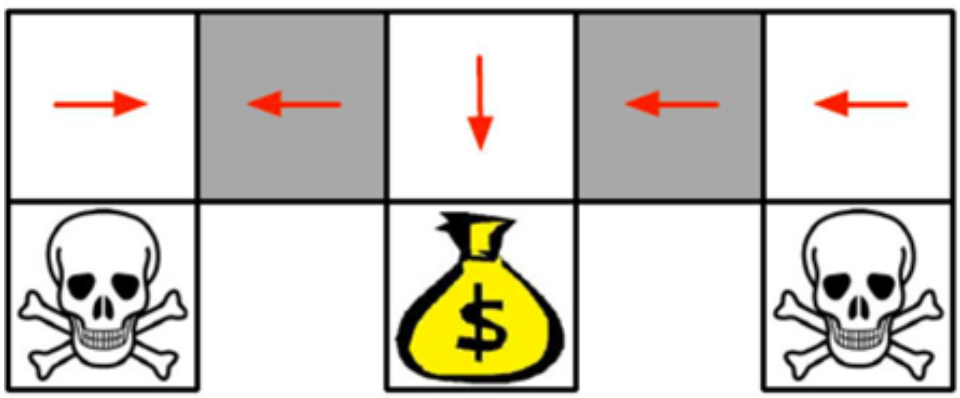

Figure 1: probabilistic policy

Consider, for example, a situation in which the Agent is able to observe that it is in one of the grey squares, but is not able to determine which grey square it is in. In this case, the policy of moving left or right with equal probability from the grey square will perform better than any deterministic policy.

In two-player games like rock-paper-scissors, a random strategy is also required in order to make agent choices unpredictable to the opponent.

2.6 Optional video

insert video

2.7 Models of Optimality

What exactly do we mean when we say that an RL Agent should choose a policy which maximises its future rewards?

The old saying "a fast nickel is worth a slow dime" reminds us that people often prefer to receive a smaller reward sooner instead of a larger reward later. Economists use surveys to gauge people's preferences in this regard: "Would you prefer ten dollars today, or 15 dollars next week? How about 50 dollars today compared to 100 dollars next year?" Responses to these surveys can be used to estimate a number *gamma* between 0 and 1 that is chosen such that one dollar in the current timestep is considered equally desirable with *gamma* dollars in the next timestep. In Reinforcement Learning, this number *gamma* is called the discount factor and is used to define the infinite discounted reward, which can be compared with average reward and finite horizon reward.

$$\text{Finite horizon reward}: \sum_{i=0}^{h-1} r_{t+1}$$

$$\text{Infinite discounted reward}: \sum_{i=0}^{\infty} \gamma^i r_{t+i}, 0 \leq \gamma \leq 1$$

$$\text{Average Reward}: \lim_{h \to \infty} \frac{1}{h} \sum_{i=0}^{h-1} r_{t+1}$$

We normally try to maximise the infinite discounted reward, because it is easier to analyse theoretically and also works well in practice.

The finite horizon reward is easy to compute, but may lead to bad decisions by failing to take into account rewards which are just over the horizon.

Average reward is hard to deal with because we cannot sensibly choose between a small reward soon and a large reward very far in the future - for example, 100 dollars today compared with a million dollars, 50 years from now.

2.8 Comparing Models of Optimality

This environment illustrates how the choice of action may depend on the model of optimality. An agent trying to maximise the finite horizon reword with $h = 4$ would prefer action a_1 because it is the only action which will gain any reward within the first four time steps. An agent maximising average reward will prefer action a_3 because, after the first few time steps, it will be

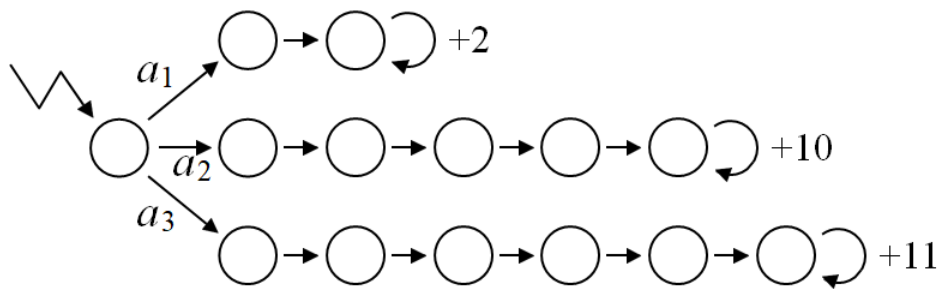

Figure 2: comparing models of optimality

getting a reward of +11 every timestep which is larger than +10 for an agent who chose action a_2. An agent maximising infinite discounted reward with $\gamma = 0.9$ will prefer action a_2. One way to see this is as follows: Consider Agent A_2 choosing action a_2, compared to Agent A_3 choosing action a_3. The reward of 11 received by A_3 in timestep 7 would be equivalent to 9.9 in timestep 6 and therefore slightly less than the \$0 received by A_2 in timestep 6. By the same logic, the \$0 received by A_2 in timestep nnn is always slightly more valuable than the \$1 received by A_3 in timestep $n + 1$.

2.9 Optional Video

insert video

2.10 Value Function

Every policy π determines a Value Function $V^\pi : S \to R$ where $V^\pi(s)$ is the average discounted reward received by an agent who begins in state sss and chooses its actions according to policy π.

2.11 video: summing an infinite series

If $\gamma = 0.9$, the value of the last node in the middle row of the above diagram is $\frac{10}{1-\gamma} = 100$. The value of the last node in the lowest row is $\frac{11}{1-\gamma} = 110$. The value of the second last node in the lower row is 99, which can be obtained from the value of the last node either by subtracting 11 or multiplying by γ.

2.12 Value Function Example

This image shows the value function V_π in a grid world environment, where π is the policy of choosing between available actions uniformly randomly.

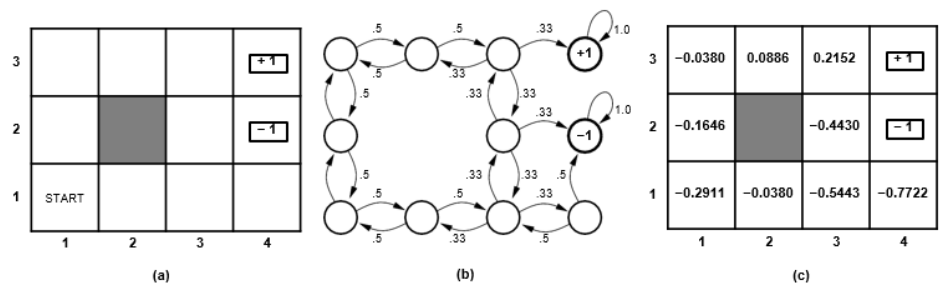

Figure 3: value function example

3 Exploration and Delayed Reinforcement

For small environments, we can compute the value function using simultaneous equations. For larger environments, we need a way of learning it, through experience.

3.1 Exploration and delayed reinforcement

I was born to try ... But you've got to make choices Be wrong or right Sometimes you've got to sacrifice the things you like. - Delta Goodrem

3.2 Multi-Armed Bandit Problem

Figure 4: multiarmbandit

The special case of a stochastic reinforcement learning environment with only one state is called the Multi-Armed Bandit Problem, because it is like being in a room with several (friendly) slot machines (also called "one-armed bandits") for a limited time, and trying to collect as much money as possible. We assume that each slot machine provides rewards chosen from its own (stationary) distribution, and that in each time step we are able to pull the lever on only one machine.

3.3 Exploration / Exploitation Tradeoff

After pulling the levers a few times and observing the rewards received, it makes sense that most of the time we should choose the lever (action) which we think will give the highest expected reward, based on previous observations.

However, in order to ensure convergence to the optimal strategy, we must occasionally choose something different from our preferred action. Perhaps the simplest way to achieve this is known as an epsilon-greedy strategy, where with probability $1-\epsilon$ we choose what we think is the best action, and with probability ϵ (typically, 5 percent) we choose a random action.

$$P(a) = \frac{e^{R(a)/T}}{\sum_{b \in A} R(b)/T}$$

More sophisticated strategies also exist such as Thompson Sampling, Upper Confidence Bound (UCB) algorithms, or choosing from a Softmax (Boltzmann) distribution based on the average reward so far observed for each action.

You may have noticed that we make the same kind of judgements about exploration versus exploitation in everyday situations. Should you eat at the old restaurant you have visited many times, or should you try the newly opened restaurant across the street which might be worse but might be much better?

3.4 Optional video

insert video

3.5 Delayed Reinforcement

The need for exploration also applies in the general case where there are multiple states, and where we may need to take a whole sequence of actions in order to get to the state from which we can obtain a high reward.

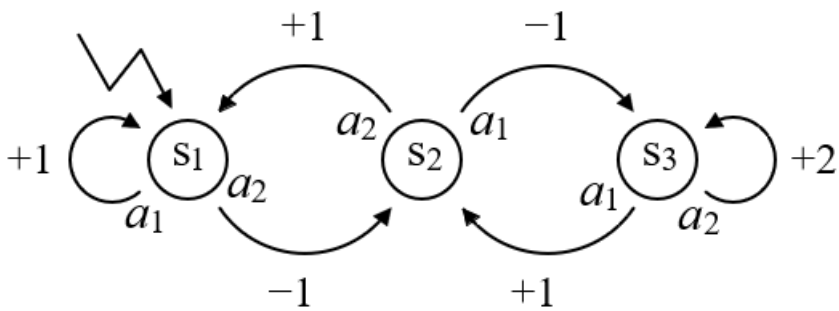

Figure 5: asdf

Recall that every policy π determines a Value Function arrow $V^\pi : S \to R$ where $V^\pi(s)$ is the expected discounted reward received by an agent who begins in state sss and chooses its actions according to policy π.

If $\pi = \pi^*$ is optimal, then $V^*(s) = V^{\pi^*}(s)$ is the maximum (expected) discounted reward obtainable from state s. Knowing this optimal value function can help to determine the optimal policy.

3.6 Computing the Value Function

Let $\gamma = 0.9$ and consider the policy $\pi : S_1 \mapsto a_2, S_2 \mapsto a_1, S_3 \mapsto a_2$. We have

$$\begin{aligned}
V^\pi(S_3) &= \frac{2}{1-\gamma} = 20 \\
V^\pi(S_2) &= -1 + \gamma V(S_3) = -1 + 18 = 17 \\
V^\pi(S_1) &= -1 + \gamma V(S_2) = -1 + 15.3 = 14.3
\end{aligned}$$

For this example, the policy π shown above must be the optimal policy, because the value function for states S_1 and S_2 under any other policy would not be more than 101010. Therefore, the optimal value function is

$$V^* : S_1 14.3, S_2 \mapsto 20$$

If we were given this value function V^*, we could use it to determine the optimal policy π provided we also know the reward function $R : S \times Z \to R$ and the transfer function $\delta : S \times A \to S$. This information is sometimes called the "World Model".

3.7 Video: Computing the Value Function

insert video

3.8 Q-Function

The Q-function is a more sophisticated version of the value function which enables an agent to act optimally without needing to know the World Model. For any policy π, the Q-function $Q^{\pi}(s, a)$ is the expected discounted reward received by an agent who begins in state s, first performs action a and then follows policy π for all subsequent timesteps.

If $\pi = \pi^*$ is optimal, then $Q^*(s, a) = Q^{\pi^*}(s, a)$ is the maximum (expected) discounted reward obtainable from sss, if the agent is forced to take action aaa in the first timestep but can act optimally thereafter.

If the optimal Q-function Q^* is known, then the optimal policy is given by

$$\pi^*(s) = \arg\max Q^*(s, a)$$

The Q-function for the environment shown above can be computed as follows:

$$\begin{aligned}
&Q(S_1, a_1) = 1 + \gamma V(S_1) = 1 + 0.9 \times 14.3 = 13.87 \\
&Q(S_1, a_2) = V(S_1) = 14.3 \\
&Q(S_2, a_1) = V(S_2) = 17 \\
&Q(S_2, a_2) = 1 + \gamma V(S_1) = 13.87 \\
&Q(S_3, a_1) = 1 + \gamma V(S_2) = 1 + 0.9 \times 17 = 16.3 \\
&Q(S_3, a_2) = V(S_3) = 20
\end{aligned}$$

3.9 Optional video

insert video

4 TD-learning and Q-learning

Reinforcement Learning algorithms can generally be grouped into three classes:

- Value function learning, including TD-Learning and Q-Learning
- Policy learning, including policy gradients and evolution strategies

- Actor-Critic, which is a combination of value function and policy learning

4.1 Value Function Learning

Recall that if $\pi = \pi^*$ is optimal, then $V^*(s) = V^{\pi^*}(s)$ is the maximum (expected) discounted reward obtainable from state s, and $Q^*(s,a) = Q^{\pi^*}(s,a)$ is the maximum (expected) discounted reward obtainable from s, if the agent is forced to take action aaa in the first timestep but can act optimally thereafter.

The idea behind value function learning is that the agent retains its own estimate $V^*()$ or $Q^*()$ of the "true" value function $V^*()$ or $Q^*()$. This estimate might be quite inaccurate to start with, but it gets iteratively improved over time so that it more closely approximates the true value. This process is sometimes called "Bootstrapping".

4.2 Temporal Difference Learning

Let's first assume that R and δ are deterministic. Then the (true) value $V*(s)$ of the current state s should be equal to the immediate reward plus the discounted value of the next state

$$V^*(s) = R(s,a) + \gamma V^*(\delta(s,a))$$

We can turn this into an update rule for the estimated value:

$$V(s_t) \leftarrow r_t + \gamma V(s_{t+1})$$

If R and $\delta arestochastic(multi-valued), itisnotsafetosimplyreplace$V(s)$withtheexpressionontheri$

$$V(s_t) \leftarrow V(s_t) + \eta[r_t + \gamma V(s_{t+1}) - V(s_t)]$$

4.3 Q-Learning

Q-learning is similar to TD-learning except that the Q-function $Q^* : S \times A \rightarrow R$ depends on a state, action pair instead of just a state.

For a deterministic environment, π^*, Q^* and V^* are related by

$$\begin{aligned}
\pi^*(s) &= \arg\max Q^*(s,a) \\
Q^*(s,a) &= R(s,a) + \gamma V^*(\delta(s,a)) \\
V^*(s) &= \max Q^*(s,b) \\
\text{so} Q*(s,a) &= R(s,a) + \gamma \max Q^*(s,b)
\end{aligned}$$

The last equation suggests that we can iteratively update our estimate by

$$Q(s_t, a_t)r_t + \gamma \max Q(s_{t+1}, b)$$

If the environment is stochastic, we instead write

$$Q(s_t, a_t)Q(s_t, a_t) + \eta[r_t + \gamma \max Q(s_{t+1}, b) - Q(s_t, a_t)]$$

4.4 Optional video

insert video

4.5 Theoretical Results and Generalization

It can be proved that TD-learning and Q-learning will eventually converge to the optimal policy, for any deterministic Markov decision process, assuming an appropriately randomised strategy (Sutton, 1988; Watkins & Dayan 1992; Dayan & S ejnowski, 1994).

However, these results rely on the assumption that every state will be visited many times. For small environments with a limited number of states and actions, this assumption is reasonable. For more challenging environments, the number of states can be exponentially large, and it may be impractical to assume that every state will be visited even once. In these situations, we need to rely on a model such as a neural network that is able to generalise to estimate the value of new states based on their similarity with states that have previously been encountered.

4.6 Optional video

insert video

4.7 Computer Game Playing

Suppose we want to write a computer program to play a game like backgammon, chess, checkers or Go. This can be done using a tree search algorithm (expectimax, MCTS, or minimax with alpha-beta pruning). But we need:

(a) an scheme for encoding any board position as a set of numbers, and (b) a way to train a neural network or other learning system to compute a board evaluation, based on those numbers.

4.8 Backgammon

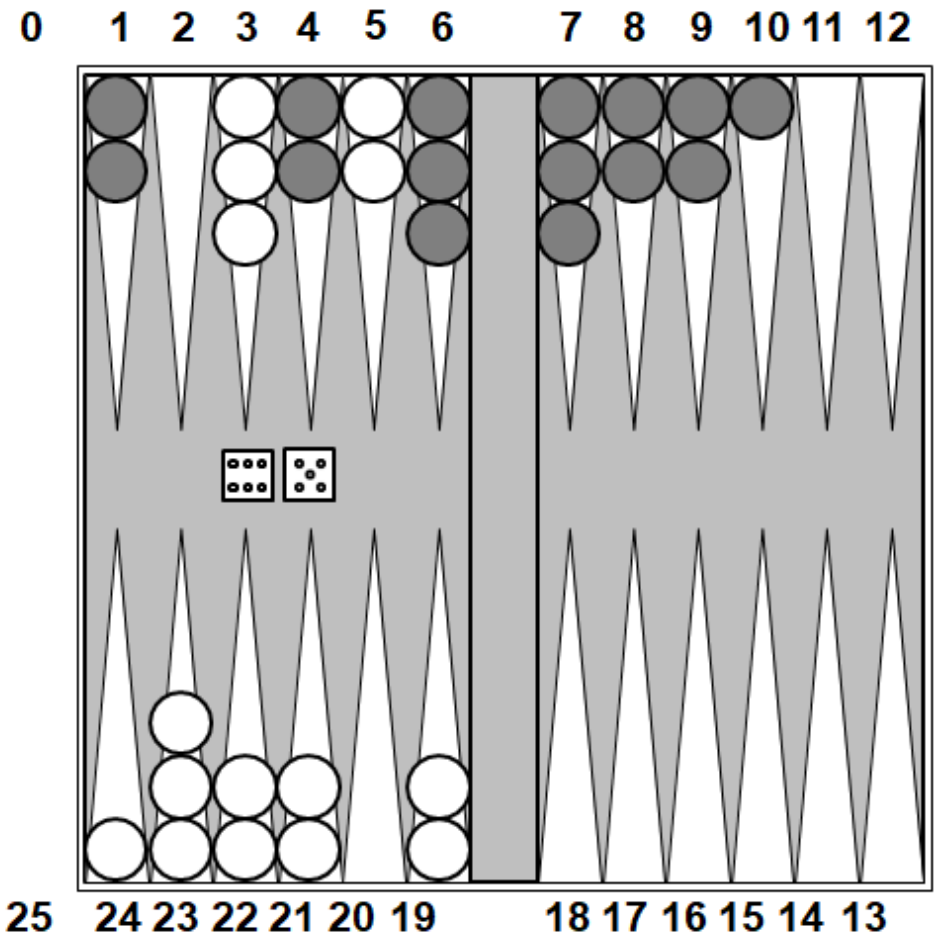

Figure 6: backgammon

One of the first practical applications of TD-learning (Tesauro, 1992) was to play the game of backgammon, which involves rolling dice and moving pieces around a board with 24 "points". In the above diagram, the white player is moving their pieces clockwise toward the bottom left quadrant while the black player is moving their pieces anti-clockwise toward the upper left quadrant. When a player gets all of their pieces into the last quadrant, they can use the dice rolls to move pieces off the board. The first player to get all of their pieces off the board is the winner.

If a player has only one piece on a point (such as point 10 or 24 above) and the other player moves a piece to that point, the original piece is sent back to the "bar" and must start again from the beginning. If a player has two or more pieces on a point, the other player is not allowed to move a piece to that point. We see in the above position that the black player has multiple pieces on points 1, 4, 6, 7, 8, 9 so that the white player will not be able to move their pieces on points 3 or 5 unless they roll either 2, 5 or 6.

4.9 TD-Learning and Q-Learning

One of the first practical applications of TD-learning (Tesauro, 1992) was to play the game of backgammon, which involves rolling dice and moving pieces around a board with 24 "points". In the above diagram, the white player is moving their pieces clockwise toward the bottom left quadrant while the black player is moving their pieces anti-clockwise toward the upper left quadrant. When a player gets all of their pieces into the last quadrant, they can use the dice rolls to move pieces off the board. The first player to get all of their pieces off the board is the winner.

If a player has only one piece on a point (such as point 10 or 24 above) and the other player moves a piece to that point, the original piece is sent back to the "bar" and must start again from the beginning. If a player has two or more pieces on a point, the other player is not allowed to move a piece to that point. We see in the above position that the black player has multiple pieces on points 1, 4, 6, 7, 8, 9 so that the white player will not be able to move their pieces on points 3 or 5 unless they roll either 2, 5 or 6.

4.10 TD-Gammon Neural Network

TD-Gammon used a 2-layer neural network with 196 inputs, 20 hidden nodes and 1 output. The inputs are arranged as follows:

- 4 units $\times$ 2 players $\times$ 24 points
- 2 units for pieces on the bar
- 2 units for pieces off the board

Four inputs are assigned to each colour and each point, with a 1-hot encoding used to specify one, two or three pieces and the fourth input scaled in proportion to the number of pieces if there are more than three. Additional inputs specify the number of pieces of each colour which are on the bar, or off the board.

For any encoded board state, the network produces an output V between 0 and 1 which is interpreted as its estimate of the probability of winning from that state.

At each timestep, the dice are rolled, the network considers all possible legal moves and computes the "next board position" that would result from each move. These are converted to the appropriate input format, fed to the network, and the one which produces the largest output is chosen.

The network can be trained by backpropagation using this formula

$$w_i \leftarrow w_i + \eta(T - V)\frac{\partial V}{\partial w_i}$$

where are the weights in the network,

$$isthelearningrate,$$

$$\mathrm{V}istheactualoutputofthenetwork, andTTTisthetargetvalue.$$

The question is: How do we determine the target value T? In other words, how do we know what the value of the current position "should" have been? Alternatively, how do we find a better estimate for the value of the current position?

4.11 How to choose the Target value

One approach is to have a human expert play many games and build a database of positions, dice rolls and chosen moves, similar to what was done for the ALVINN autonomous driving system we discussed in Week 1. The network could then be trained to increase the evaluation of each move chosen by the human expert and decrease that of other moves which were available but did not get chosen.

Another approach is for the network to learn from self-play by TD-learning, effectively using the evaluation of subsequent positions in the game to refine the evaluation of earlier positions.

Other methods such as TD-Root, TD-Leaf, MCTS and TreeStrap (Veness et al., 2009) combine learning with tree search. These are important for deterministic games like chess or Go, but less important for a stochastic game like backgammon because the randomness of the dice rolls limits the benefit of deep lookahead.

For backgammon, the agent receives just a single reward at the end of each game, which we can consider as the final value (typically, +1 for a win or -1 for a loss). We then have a sequence of game positions, each with its own (estimated) value:

$$\text{(current estimate)}V_t \to V_{t+1} \to \cdots \to V_m \to V_{m+1}\text{(final result)}$$

In this context, TD-learning simplifies and becomes equivalent to using the value of the next state as the training value for the current state. A fancier version, called TD, uses a weighted average over future estimates as the training value for, where

$$T_k = (1-\lambda)\sum_{k=t+1}^{m} \lambda^{k-1-t} V_k + \lambda^{m-t} V_{m+1}$$

The parameter λ between 0 and 1 serves as a kind of discount factor, but is different to the usual discount factor γ (which can be equal to 1 in the case of TD-Gammon, because we do not care when we get the reward so long as we get it).

4.12 TD-Gammon

Tesauro trained two networks - one using Expert Preferences (EP), the other by TD-learning (TD). The TD-network outperformed the EP-network, and with modifications such as 3-step lookahead (expectimax) and additional hand-crafted input features, became the best backgammon player in the world in 1995.

4.13 Optional video

insert video

4.14 Exercise: Reinforcement learning

Consider an environment with two states $S = \{ S_1, S_2\}$ and two actions $A = \{ a_1, a_2\}$ where the deterministic transitions δ and reward R for each state and action are as follows:

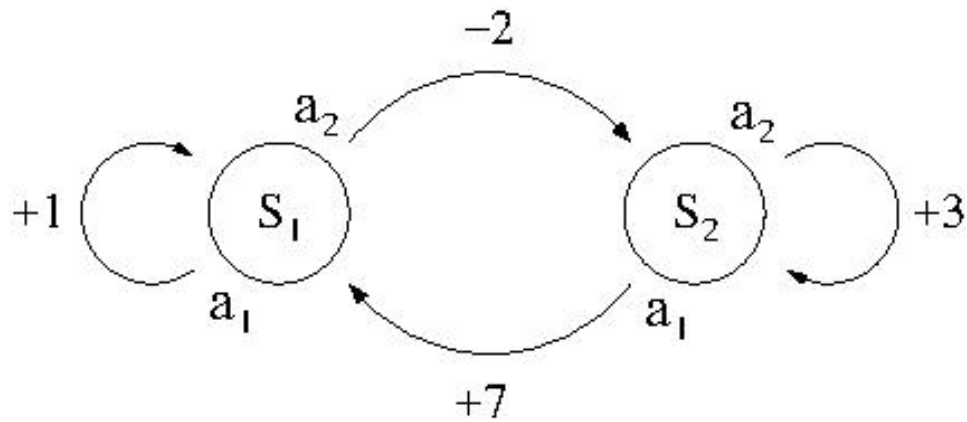

Figure 7: question 1

$$\delta(S_1, a_1) = s_1, R(S_1, a_1) = +1$$
$$\delta(S_1, a_2) = s_2, R(S_1, a_2) = -2$$
$$\delta(S_2, a_1) = s_1, R(S_2, a_1) = +7$$
$$\delta(S_2, a_2) = s_2, R(S_2, a_2) = +3$$

question 1 Assuming a discount factor of $\gamma = 0.7$, determine the optimal policy $\pi^* : S \to A$

Answer

$$\pi^*(S_1) = a_2$$
$$\pi^*(S_2) = a_1$$

Question 2 Still assuming $\gamma = 0.7$ determine the value function $V : S\ toR$

Answer

$$V(S_1) = 5.69$$
$$V(S_2) = 10.98$$

Explanation

$$V(S_1) = -2 + \gamma V(S_2)$$
$$V(S_2) = +7 + \gamma V(S_1)$$
$$\text{So} V(S_1) = -2 + 7\gamma + \gamma^2 V(S_1)$$
$$\text{i.e.} V(S_1) = \frac{-2 + 7\gamma}{1 - \gamma^2} = \frac{-2 + 7 \times 0.7}{1 - 0.49} = 5.69$$
$$\text{Then} V(S_2) = 7 + 0.7 \times 5.69 = 10.98$$

Question 3 still assuming $\gamma = 0.7$, determine the values of the Q-function $Q : S \times A \rightarrow R$

Answer

$$\begin{aligned} Q(S_1, a_1) &= 1 + \gamma V(S_1) &&= 4.98 \\ Q(S_1, a_2) &= V(S_1) &&= 5.69 \\ Q(S_2, a_1) &= V(S_2) &&= 10.98 \\ Q(S_2, a_2) &= 3 + \gamma V(S_2) &&= 10.69 \end{aligned}$$

Question 4 Still assuming $\gamma = 0.7$, trace through the first few steps of the Q-learning algorithm, assuming a learning rate of 1 and with all Q values initially set to zero. Explain why it is necessary to force exploration through probabilistic choice of actions, in order to ensure convergence to the true Q values.

Here are some hints to get you started:

Since the learning rate is 1 (and the environment deterministic) we can use this Q-Learning update rule:

$$Q(S, a) \leftarrow r(S, a) + \gamma \max Q(\delta(S, a), b)$$

Let's assume the agent starts in state S_1. Because the initial Q values are all zero, the first action must be chosen randomly. If action is chosen, the agent will get a reward of +1 and the update will be

$$Q(S_1, a_1) \leftarrow 1 + \gamma \times 0 = 1$$

Answer With a deterministic environment and a learning rate of 1, the Q-learning update rule is

$$Q(S, a) \leftarrow r(S, a) + \gamma \max Q(\delta(S, a), b)$$

if the agent starts in a state S_1 and chooses action a_1, it will get a reward of +1 and the update will be:

$$Q(S_1, a_1) \leftarrow 1 + \gamma \times 0 = 1$$

we do **not** force exploration, the agent will always prefer action a_1 in state S_1, and will never explore action a_2 this means that $Q(S_1, a_2)$ will remain zero forever, instead of converging to the true value of 5.69. If we **do** force exploration, the next steps may look like this:

current state	chosen state	new Q value
$cell4S_1$	$cell5a_2$	$-2 + \gamma\times = -2$
S_2	a_2	$3 + \gamma \times 0 = 3$

At this point the table looks like this:

Q	a_1	a_2
S_1	1	-2
$cell7S_2$	0	3

Again we need to force exploration in order to the get the agent to choose a_1 from S_2 and to again choose a_2 from S_1

current state	chosen action	new Q value
$cell4S_2$	a_1	$7 + \gamma \times 1 = 7.7$
S_1	a_2	$-2 + \gamma 7.7 = 3.39$

Q	a_1	a_2
S_1	1	3.39
S_2	7.7	3

Further steps will refine the Q value estimates, and in the limit they will converge to their true values

Question 5 Now let's consider how the value function changes as the discount factor γ varies between 0 and 1. There are four deterministic policies for this environment, which can be written as $\pi_{11}, \pi_{12}, \pi_{21}$, and π_{22}, where $\pi_{ij}(S_1) = a_1, \pi_{ij}(S_2) = a_j$

Calculate the value function $V^{\pi}_{(\gamma)} : S \to R$ for each of these four policies (keeping γ as a variable)

Answer

Question 6 Determine for which range of values γ each of the policies $\pi_{11}, \pi_{12}, \pi_{21}, \pi_{22}$ is optimal

Answer π_{11} is optimal when:

$$0 < V^{\pi}_{11}(S_1) - V^{\pi}_{21}(S_1) = \frac{(1+\gamma) - (-2+7\gamma)}{1-\gamma^2} = \frac{3-6\gamma}{1-\gamma^2} \text{i.e.} 0 \leq \gamma \leq 0.5$$

π_{22} is optimal when

$$0 < V^{\pi}_{22}(S_2) - V^{\pi}_{21}(S_2) = \frac{3(1+\gamma) - (7-2\gamma)}{1-\gamma^2} = \frac{-4+5\gamma}{1-\gamma^2} \text{i.e.} 0.8 \leq \gamma \leq 1.0$$

π_{21} is optimal for $0.5 \leq \gamma \leq 0.8$

π_{12} is never optimal because it is dominated by π_{11} when $\gamma < 2/3$ and by π_{22} when $\gamma > 0.6$

$$
\begin{aligned}
V^{\pi}_{11}(S_1) &= 1 + \gamma V^{\pi}_{11}(S_1), \quad \text{so} \quad V^{\pi}_{11}(S_1) = \frac{1}{1-\gamma} \\
V^{\pi}_{11}(S_2) &= 7 + \gamma V^{\pi}_{11}(S_1) = 7 + \frac{\gamma}{1-\gamma} \\
V^{\pi}_{12}(S_1) &= V^{\pi}_{11}(S_1) = \frac{1}{1-\gamma} \\
V^{\pi}_{12}(S_2) &= \frac{3}{1-\gamma} \\
V^{\pi}_{21}(S_1) &= -2 + 7\gamma + \gamma^2 V^{\pi}_{21}(S_1), \quad \text{so} \quad V^{\pi}_{21}(S_1) = \frac{-2+7\gamma}{1-\gamma^2} \\
V^{\pi}_{21}(S_2) &= 7 - 2\gamma + \gamma^2 V^{\pi}_{21}(S_2), \quad \text{so} \quad V^{\pi}_{21}(S_2) = \frac{7-2\gamma}{1-\gamma^2} \\
V^{\pi}_{22}(S_1) &= -2 + \frac{3\gamma}{1-\gamma} \\
V^{\pi}_{22}(S_2) &= \frac{3}{1-\gamma}
\end{aligned}
$$

Figure 8: 5a_5

4.15 Exercise: Reinforcement Learning

This is a revision quiz to test your understanding of the material from Week 5 on Reinforcement Learning.

You must attempt to answer each question yourself, before looking at the sample answer.

Question 1 Describe the elements (sets and functions) that are needed to give a formal description of a reinforcement learning environment. What is the difference between a deterministic environment and a stochastic environment?

Answer Formally, a reinforcement learning environment is defined by a set of $\mathcal{S}$ states, a set of $\mathcal{A}$, a transition function δ and a reward function $\mathcal{R}$

For a stochastic environment δ and $\mathcal{R}$ are single-valued functions:

$$
\begin{aligned}
\delta &: \mathcal{S} \times \mathcal{A} \to \mathcal{S} \\
\mathcal{R} &: \mathcal{S} \times \mathcal{A} \to \mathbb{R}
\end{aligned}
$$

For a stochastic environment δ and/or $\mathcal{R}$ are not single-valued, but instead define a probability distribution on $\mathcal{S}$ or $\mathbb{R}$

Question 2 Name three different models of optimality in reinforcement learning, and give a formula for calculating each one.

Answer

$$\text{Finite horizon reward}: \sum_{i=0}^{h-1} r_{t+i}$$

$$\text{Infinite discounted reward}: \sum_{i=0}^{\infty} \gamma^i r_{t+i}, 0 \leq \gamma < 1$$

$$\text{Average reward}: \lim_{h\to\infty} \frac{1}{h} \sum_{h-1}^{i=0} r_{t+i}$$

Question 3 What is the definition of:

- the optimal policy
- the value function
- the Q-function?

Answer the optimal policy is the function $\pi^* : \mathcal{S} \to \mathcal{A}$ which maximises the infinite discounted reward.

the value function $V^\pi(S)$ is the expected infinite discounted reward received by an agent who begins in state s, first performs action a and then $V^*(s) = V^{\pi^*}(s)$ is the maximum (expected) infinite discounted reward obtainable from state s

The Q-function $Q^\pi(spa)$ is the expected infinite discounted reward received by an agent who begins in state s, first performs action a and then follows policy π for all subsequent timesteps If $\pi = \pi^*$ is optimal, then $Q^*(s,a) = Q^{\pi*}(s,a)$ is the maximum (expected) discounted reward obtainable from s if the agent is forced to take action a in the first time step but can act optimally thereafter.

Question 4 Assuming a stochastic environment, discount factor γ and learning rate of η, write the equation for

- Temporal Difference learning TD(0)
- Q-Learning

Remember to define any symbols you use.

Answer

$$V(s_t) \leftarrow V(s_t) + \eta[r_t + \gamma V(s_{t+1}) - V(s_t)]$$
$$Q(s_t, a_t) \leftarrow Q(s_t a_t) + \eta[r_t + \gamma \max_b Q(s_{t+1}, b) - Q(s_t, a_t)]$$
$$s_t = \text{ state at time } t, a_t = \text{ action performed at time} t,$$
$$r_t = \text{ reward received at time } t, s_{t+1} = \text{ state at time } t+1$$

Part II
Policy Learning and Deep RL

5 Policy Learning

Policy learning algorithms do not use a value function but instead operate directly on the policy, chosen from a family of policies $\pi_\theta : SA$ determined by parameters θ.

Typically, π_θ is a neural network with weights θ which takes a state s as input and produces action a as output, which may be either continuous or discrete.

If there is a discrete choice of actions, the network has one output for each possible action and uses Softmax to determine the conditional probability of performing action a in state s.

For episodic domains like Backgammon, we do not need a discount factor, and the "fitness" of policy π_θ can be taken as the Value function of the initial states s_0 under this policy, which is the expected (or average) total reward received in each game by an agent using policy π_θ.

$$\text{fitness}(\pi_\theta) = V^{\pi\theta}(s_0) = E_{\pi\theta}(r_{\text{total}})$$

Policy Learning algorithms include Policy Gradients, which use gradient descent to modify the parameters

$, andEvolutionStrategies, whichmakerandomchangestoandkeeponlythoseupdatesthatareseentoincre$

6 Deep Reinforcement Learning

6.1 Deep Reinforcement Learning

6.2 Deep Q-Learning for Atari Games

Mnih (2015) demonstrated how Q-learning could be combined with deep CNNs to learn to play Atari games from pixels.

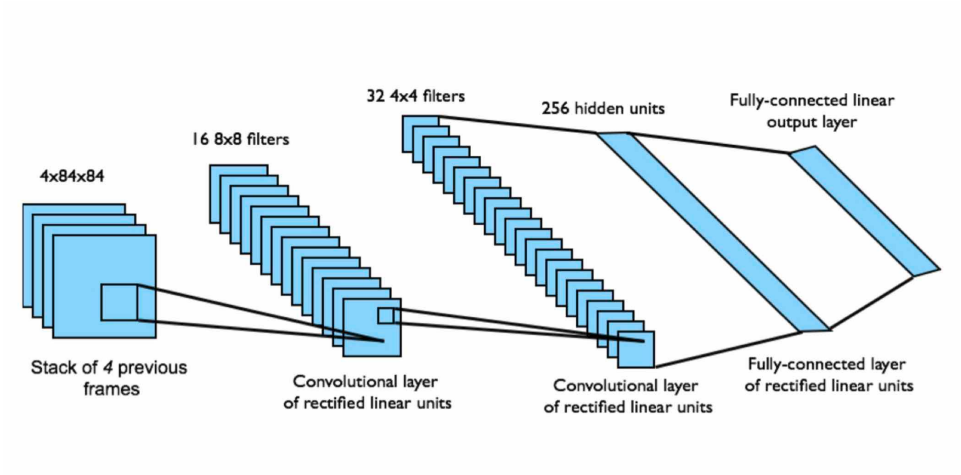

Figure 9: atari games

The input state s is a stack of raw pixels from four consecutive frames. The images are converted from 8-bit RGB images at resolution 210×160 pixels to 84×84 greyscale images. The 18 outputs are the Q-values $Q(s, a)$ for 18 different combinations of joystick/button positions. The reward is the change in score during the timestep.

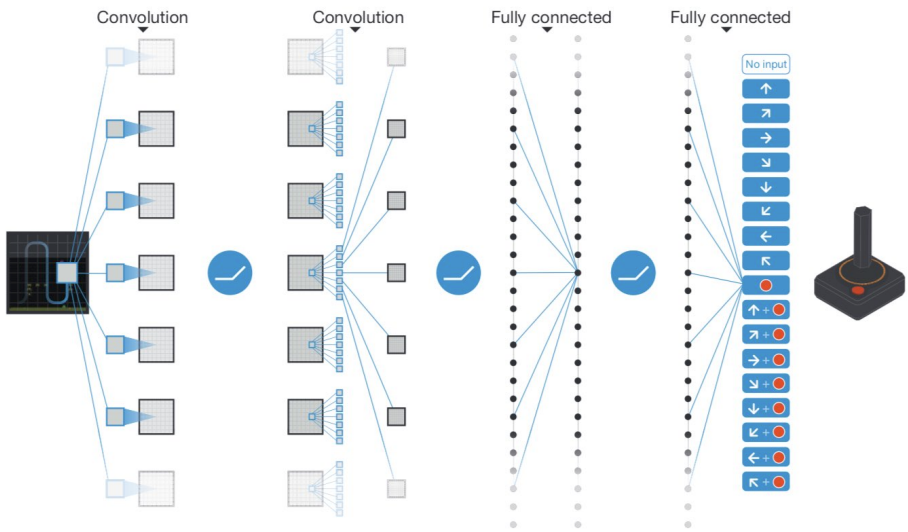

Figure 10: deep q learning

Recall the Q-Learning update rule:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \eta \Big[r_t + \gamma \max Q(s_{t+1}, b) - Q(s_t, a_t) \Big]$$

If a lookup table is used to store the Q values for every state and action separately, the algorithm is guaranteed to eventually converge to an optimal policy. But, if the number of states is exponentially large, we must instead use a neural network Q_w and adjust the weights www according to the Q-Learning update rule, which is equivalent to minimizing:

$$\Big[r_t + \gamma \max Q_w(s_{t+1}, b) - Q_w(s_i, a_t) \Big]^2$$

The gradient is applied only to $Q_w(s_t, a_t)$, not to $Q_w(s_{t+1}, b)$.

7 Experience Replay

Training of deep neural networks for classification tasks generally requires that each mini-batch should include a variety of different inputs and target outputs. For Atari games, many similar states may occur in succession,

often with the same action being selected. We can remove this **temporal** correlation between samples by storing experiences in a Replay Buffer.

In this scenario, one thread repeatedly plays the game, selecting its actions using an ϵ-greedy strategy based on the current ϵ-values, and builds a database of experiences s_t, a_t, r_t, s_{t+1}. Another thread samples asynchronously from this database and applies the Q-learning rule to minimize

$$\left[r_t + \gamma \max Q_w(s_{t+1}, b) - Q_w(s_t, a_t)\right]^2$$

This removes temporal correlations by sampling from a variety of game situations in random order, and also makes it easier to parallelize the algorithm on multiple GPUs.

8 Optional video

9 Prioritised Replay

Instead of sampling experiences uniformly, it may be more efficient to store and retrieve them in a priority queue with priority based on the DQN error (Schaul, 2015).

$$|r_t + \gamma \max Qw(s_{t+1}, b) - Q_w(s_t, a_t)|$$

This ensures the system will concentrate more effort on situations where the value was "surprising" (in the sense of being far away from what was predicted).

10 Double Q-Learning

If the same weights www are used to select actions and to evaluate actions (as well as states) the network may learn a suboptimal strategy due to a kind of "confirmation bias". One way to avoid this is to maintain two sets of weights www and , with one used for action selection and the other for evaluation (then swap their roles).

In the context of Deep Q-learning, a simpler approach is to use the current "online" version of www for selection, and an older "target" version for evaluation (Van Hasselt, 2016). We therefore minimize

A new version of is periodically calculated from the distributed values of

$$andthis$$

is broadcast to all processors.

11 Advantage Function

The Function can be written as a sum of the value function plus an advantage function represents the advantage (or disadvantage) of taking action aaa in state s, compared to taking the action preferred by the current policy . We can learn approximations for these two components separately:

Note that actions can be selected just using , because

12 Optional video

12.1 Advantage Actor-Critic

Recall that in the REINFORCE algorithm, a baseline bbb could be subtracted from for the purpose of variance reduction.

In the actor-critic framework,

We can also subtract a baseline from . This baseline must be independent of the action at, , but it could be dependent on the state . A good choice of baseline is the value function , in which case the Q function is replaced by the advantage function

12.2 Asynchronous Advantage Actor-Critic

The Asynchronous Advantage Actor-Critic or Algorithm combines a policy network , a value function network and an (estimated) Q-function.

- use polity network to choose actions
- learn a parameterised value function (s) by TD-learning
- estimate Q-value by n-step sample
- update policy by
- update value function by minimising

12.3 Optional video

13 Examples of Deep Reinforcement Learning